

A Computer-Vision-Based Bikeability Score

Bünyamin Aydemir, Gonzalo Guerrero Salazar, and Hendrick Fischer

Abstract—

Index Terms—Bikeability, computer vision, zero-shot classification, vision-language models, CLIP, surface classification, object detection, GPX.

I. INTRODUCTION

Computer vision is one of the technologies that has firmly established itself in recent decades. The increase in computing capacity, the abundance of data, and the improvement of algorithms have propelled it—just as they have other areas of artificial intelligence—to the point where the technology is becoming a standard tool. We now view it as the de-facto solution to problems that previously required entirely different methods.

Take ADAS, for example, which are now expected in passenger vehicles rather than viewed as luxury equipment. However, these systems have largely not made their way to smaller vehicles like motorcycles or bicycles, the latter of which provide transportation and recreation to millions of people. Regular bicycle commuting inspired this project to implement vision systems not only to assess the safety, but also the aesthetic—and therefore mood-improving—aspects of bicycle routes. We call this metric the “bikeability score”, based on the premise that while all roads are technically drivable, they do not all provide the same riding experience.

The score singles out three distinct factors that define this experience:

- **The road itself (Surface):** Determining the type of surface the user is riding on to assess if the vehicle is adequate and the ride will be physically comfortable.
- **The obstacles along the route (Traffic Dynamics):** Evaluating whether traffic is dominated by other bicycles or cars, and identifying major crossings to determine if the ride can remain largely continuous.
- **The surroundings (Environment):** Recognizing whether the route is dominated by cluttered cityscapes or greenery, and detecting natural features like forests and rivers to assess if the route is enjoyable rather than just a functional commute.

These three factors are defined as our target sub-functions, as each contributes differently to the riding experience and operates under distinct dynamics. Consequently, each task required a dedicated approach, making it essential to select models specifically suited to each sub-function. Similarly, dataset annotation and labeling were tailored to the specific demands of each individual task. The following sections explain the methodology for selecting, analyzing, and adapting

these individual models. The final section details how the overall score is calculated from the combined outputs of each pipeline.

II. RESEARCH APPROACH AND METHODOLOGY

The system is implemented as a multi-stream pipeline. Video data, collected using a handlebar-mounted DJI camera, is sampled at a frequency of one frame every five seconds. This dataset is passed in batches of eight to the individual detector functions. The resulting feature arrays are then processed by the primary scoring function, with each frame’s evaluated score appended to an output file representing the main product of this analysis. Downstream from this output, several interpretation methods are possible, including min/max analysis and score averaging. For this project, a visualization was also generated by integrating GPS data.

The methodology initially leveraged pre-trained and zero-shot models where possible to minimize manual labeling effort and allow the taxonomy to evolve. However, as demonstrated in subsequent sections, this approach encountered specific limitations. As a result, each detector was independently benchmarked before integration into the final pipeline.

Table I gives the heuristic value vectors that turn each detector’s output into a sub-score, applied in Section VII.

Table I
HEURISTIC VALUE VECTORS FOR THE THREE SUB-SCORES ($\mathbf{q}^g$ SURFACE, $\mathbf{q}^e$ ENVIRONMENT, λ OBJECTS; NEGATIVE = BONUS).

Component	Class	Value
Ground $\mathbf{q}^g$	Cycleway	1.0
	Road	0.7
	Gravel	0.3
	Unpaved	0.1
Environment $\mathbf{q}^e$	Vegetation	1.0
	Water	1.0
	City	0.4
Object weight λ	Bicycle	−0.2
	Car	+1.0
	Traffic light	+0.3

III. OBJECTIVES AND CONSTRAINTS

The primary objective is to demonstrate a functional pipeline capable of evaluating real-world data. Consequently, testing the framework on self-collected footage was a strict requirement. However, this imposed a key constraint: ground-truth labeling had to be generated internally, either manually or with semi-automated assistance. This limitation constrained both the volume of real-world ground truth data available and the scope of benchmarking that could be conducted on candidate models.

IV. ROAD-SURFACE CLASSIFICATION

The ridden surface strongly affects cycling comfort and safety: a paved cycleway is effortless, whereas loose gravel or bare soil is strenuous and, on a road bike, hazardous. Ground detection assigns each sampled frame a single surface class and supplies the surface term of the score. Frames come from a handlebar-mounted camera at one frame every five seconds; as each is dominated by one surface, the task is single-label classification, not segmentation.

A. Task Definition and Model Selection

Two model families can solve this task: a classifier *trained on our own labels*, or a *zero-shot* vision-language model (VLM). We avoid a self-trained classifier, which would need a large, balanced, manually labeled training set and re-training on every change of the class definitions. A zero-shot VLM of the CLIP family [1] instead classifies an image by its similarity to natural-language class descriptions, so the taxonomy is adapted by editing text alone.

B. Surface Classes and Prompts

We distinguish four cycling-relevant classes: **Cycleway** (a dedicated paved bike path), **Road** (a paved road for motor traffic), **Gravel** (loose crushed stone), and **Unpaved** (soil, mud, or field tracks). Each class is described by several English prompts whose normalized embeddings are averaged (*prompt ensembling*). As Cycleway and Road share the same smooth asphalt, the prompts encode the *discriminating context* (a narrow separated bicycle lane versus a wide multi-lane road) to separate the two embeddings. A frame is assigned to the class of highest cosine similarity.

C. Evaluation Protocol

We compare nine open CLIP-family models of light to medium size through the unified `open_clip` API [2]: OpenAI CLIP, OpenCLIP (LAION-2B, DataComp-XL) [3], SigLIP [4], MobileCLIP [5], and EVA-CLIP [6]; both *accuracy* and *latency* matter for on-bike use. For the quantitative evaluation we manually labeled more than 1,700 frames from self-recorded rides, excluding frames without a clear surface. As the classes are imbalanced, we report accuracy, balanced accuracy and macro-F1 with 95% bootstrap confidence intervals and McNemar tests; latency is measured at batch size 1 with warm-up runs (ms/frame and FPS).

D. Results

Table II lists the outcome, sorted by macro-F1. The best model, **OpenCLIP ViT-B/32 (LAION-2B)**, reaches macro-F1 0.58 at only 115 ms/frame, the best accuracy-latency trade-off. Two points stand out. First, the highest plain *accuracy* (0.78, CLIP ViT-B/16) does *not* give the highest macro-F1: accuracy rewards the majority classes, whereas macro-F1 weights all four surfaces equally and is the more faithful measure on imbalanced data. Second, the 95% bootstrap accuracy intervals of the leading models overlap, so their differences are *not*

Table II
ZERO-SHOT SURFACE-CLASSIFICATION BENCHMARK ON THE SELF-LABELED TEST SET (> 1,700 FRAMES), SORTED BY MACRO-F1.
BACC: BALANCED ACCURACY; LATENCY AT BATCH SIZE 1.

Model	P[M]	Acc	BAcc	mF1	ms	FPS
OpenCLIP B/32 (LAION-2B)	151	0.75	0.72	0.58	115	8.7
SigLIP B/16 (WebLI)	203	0.75	0.62	0.57	259	3.9
CLIP B/16 (OpenAI)	150	0.78	0.62	0.57	261	3.8
MobileCLIP-S2	99	0.73	0.70	0.56	304	3.3
CLIP B/32 (OpenAI)	151	0.70	0.60	0.54	93	10.7
EVA02-B/16	150	0.76	0.62	0.54	304	3.3
OpenCLIP B/32 (DataComp-XL)	151	0.75	0.65	0.52	125	8.0
CLIP L/14 (OpenAI)	428	0.64	0.50	0.42	929	1.1

statistically significant; only CLIP ViT-L/14 (428 M) falls clearly below them; the largest model is both slowest and weakest (0.42 macro-F1), confirming that capacity alone does not help.

Broken down per class, the best model scores F1 0.80 (Cycleway) and 0.73 (Road) but only 0.45 (Gravel) and 0.33 (Unpaved); as macro-F1 averages the four equally, the two strong classes are pulled down by the two weak ones (Fig. 1). Crucially, the weakness is one of *precision*, not recall: Unpaved is still recalled well (0.77), yet its precision collapses to 0.21 because, with only 13 true frames, a few false positives bleeding in from the majority classes suffice to ruin it; Gravel (51 frames) behaves the same way, while the residual Cycleway/Road confusion (both smooth asphalt) is comparatively minor. A macro-F1 of 0.58 is thus by design strict on such tiny classes; the accuracy (0.75) and balanced accuracy (0.72, the highest of all models) are more flattering yet still fair, and for a zero-shot classifier on strongly imbalanced data the result is solid. We therefore adopt OpenCLIP ViT-B/32 (LAION-2B) as the ground detector; as all models share the same pipeline, the differences reflect architecture and training data (latencies are hardware-dependent).




Figure 1. Row-normalized confusion matrix of the best model (OpenCLIP ViT-B/32, LAION-2B) on the labeled test set. Each cell shows the count and the row share; the diagonal is correct. Errors concentrate on the Cycleway/Road pair and the rare Gravel/Unpaved classes.




Figure 2. YOLO model comparison for the object-detection module. (a) Mean inference time vs. macro MAE across YOLO variants. (b) Macro MAE vs. confidence threshold per model.

V. OBJECT DETECTION

For this evaluator, the methodology focuses on counting objects. This design choice was made because objects move relative to the bike, as the rider is in continuous motion, so an object’s size within the field of view changes on each frame and is not a stable factor for calculations. This led to another consideration for the scope of the “bikeability score” evaluator: the same route will yield different results depending on the moment the footage is captured. Therefore the goal was set, that it must be capable of running in real time on a device mounted on the bicycle. This applies only to the object detection. Hence, additional selection criteria for are the model’s inference time and size. Larger models can also extract more complex features, creating a trade-off between inference time and accuracy. Given these constraints, the YOLO family of models were the starting point for benchmarking, and several generations and sizes of models were evaluated.

A. Relevant Objects

For the calculation of the score, the methodology restricts the categories of relevant objects. Because the project employs a pre-trained YOLO model, it is efficient to leverage its baseline capabilities. Consequently, the COCO [7] dataset convention was selected as reference, which consists of 80 object categories, from which three traffic relevant subsets were taken: **motor vehicles**, **other road users**, and **traffic infrastructure**, then representative classes were selected: **car**, **bicycle**, and **traffic light**.

B. YOLO Benchmark

Benchmarking helps isolate the strongest candidate models. Trials were conducted on T4 GPUs via Google Colab, which represent a standard benchmarking platform [8] that leverages the NVIDIA ecosystem which closely mirrors the software stack of devices like the Jetson series. Comparing models reveals that model size is the primary driver of performance cost. At small scales of the benchmark, YOLO11s surprisingly achieves a lower macro Mean Absolute Error (MAE) than YOLO26s with comparable or faster inference, maintaining this margin across most confidence thresholds (Fig. 2(a)).

Larger variants offer higher overall accuracy, but double the inference time. Model scaling carries a cost: stepping from nano to medium increases parameters eightfold (~ 3.2 M to




Figure 3. Distribution of bicycle detections across the dataset. (a) Spatial density of detections in normalised image coordinates. (b) Histogram of detected bounding box heights in pixels.

~ 25.9 M) and compute nearly tenfold (8.7 to 78.9 GFLOPs). This kind of analysis is habitual in edge deployment scenarios [9], and YOLO11s remains the stronger candidate here given its maturity and efficiency. These findings reflect behavior on our labeled data; performance on other datasets may vary.

C. Effect of Confidence Thresholds on Model Performance

Across all evaluated models, the MAE drops sharply as the threshold is raised from very low values (0.05–0.15). Raising the threshold eliminates the false positives, retaining most true positives up to a point beyond from which it filters valid detections. Tested models reach their minimum MAE at different thresholds (Fig. 2(b)). This indicates different calibrations across architectures, confirming that both model selection and threshold configuration must be picked for the target application. For our data, YOLO11s achieves its best performance at approximately 0.20.

D. Dataset Distribution and Bounding Box Analysis

The distribution of detected objects was analyzed to understand the data and evaluate the suitability of the chosen model. YOLO26x was used as a teacher model to provide bounding boxes across all frames. It was shown that cars represent the most frequent object class, while bicycles are sparser and more challenging to detect. Detections cluster along the road surface, and around the horizontal center of the image (roughly $x = 0.40$ to 0.65) at mid-to-lower frame height ($y \approx 0.45$ to 0.65), exhibiting increasing density with two hotspots near the vanishing point (Fig. 3(a)).

Object size correlates with distance, making farther objects harder to detect. Bicycles represent the most difficult category due to its relative scarcity, making their size distribution particularly relevant. Across 1001 evaluated frames, 565 raw bicycle detections were recorded, with a median bounding box height of only 93 px. This indicates that most bicycle detections occur while still distant, with few occurring at close range off to the side. An initial hypothesis suggested that the MAE for bicycles suffers from undercounting small, distant objects. At its baseline setting (confidence threshold = 0.20), the model predicted 154 bicycle instances, achieving a bicycle MAE of 0.127. Applying a minimum height filter of 70 px (discarding small, distant detections) degraded performance: count dropped to 125, and the MAE increased to 0.15. These results directly support the hypothesis and hint that excessive downscaling would erroneously eliminate detections.

VI. ENVIRONMENT DETECTION: ZERO-SHOT VS. FINE-TUNED SEMANTIC SEGMENTATION

Riding through different environments shapes the experience: green *vegetation* and *water* make a ride pleasant, too much *city* and traffic make it strenuous. Environment detection is a *multi-label whole-image* task (a frame may show several classes or none) and supplies the *largest* score weight ($w_e = 0.45$, Eq. (2)).

A. Task and Model Comparison

Task: find the best model to classify the environment on our own test frames, captured with a pod-mounted camera on the bike. **Comparison:** training-free *zero-shot CLIP* [1] vs. a *fine-tuned SegFormer* [10].

CLIP (zero-shot) scores each frame against a per-class prompt pair in a shared embedding space, yielding *one present-probability per class* – no *where* or *how much* – thresholded into a label, and is tuned only via prompts and thresholds. **SegFormer (fine-tuned)** instead labels *every pixel* and reports each class’s *area fraction* (pixels ÷ total), the *how much* the score consumes; only it yields those fractions and a pixel-level mIoU, at the cost of training. Given limited time and data, SegFormer is fine-tuned on *public* Mapillary [11] masks and CLIP tuned on a *small labelled validation split* of our own frames; both are evaluated on the same own-frame test set.

B. Zero-Shot CLIP: Optimization

Two levers dominated (Table III): *per-class threshold calibration* on a leakage-free validation split (+0.095 macro-F1) and *prompt ensembling* [1] over 6–8 templates per class (+0.044); larger capacity (ViT-L/14) did *not* help. Adding the water rides in the final test (C7) then exposed CLIP’s apparent water strength as a small-validation artifact.

Table III
CLIP ZERO-SHOT OPTIMIZATION STEPS: C1–C6 ARE PRE-WATER; C7 ADDS THE WATER RIDES. FINAL: ViT-B/32, PROMPT ENSEMBLING.

#	Optimization step	Macro-F1
C1	Aligned prompt pairs, global threshold	0.547
C2	Ride-disjoint val/test split	0.632
C3	Per-class threshold calibration	0.727
C4	Prompt ensembling, 6–8 templates	0.771
C5	Larger backbone ViT-L/14 (rejected)	0.756
C6	3-class taxonomy (veg. merge)	0.846
C7	Water-inclusive final test	0.649

C. Fine-Tuned SegFormer: Optimization

Tuning ran in two phases (Table IV). (i) *Training* on Mapillary [11] pixel masks, by **mIoU**: the decisive lever was *data* – balanced oversampling of the few water images plus 258 ADE20K scenes [12] lifted water IoU $0 \rightarrow 0.70$, and an ADE20K base [13] with an “other” class raised it further. (ii) *Evaluation* on held-out own frames, by **macro-F1**, ending on the unseen water-inclusive test (0.704); ablation S10 values fine-tuning at +0.027.

Table IV
SEGFORMER OPTIMIZATION IN TWO PHASES: *training* ON MAPILLARY PIXEL MASKS (mIoU, DEV SET), THEN *evaluation* ON OWN FRAMES (MACRO-F1).

#	Optimization step	mIoU _{tr}	F1 _{eval}
S1–2	Cityscapes-B0, random data; LR tuned	0.62	–
S3–4	Balanced + water aug + ADE water	0.82	–
S5	Base swap $\rightarrow$ ADE20K (water prior)	0.857	–
S7	First own eval + val thresholds	–	0.636
S8	Explicit “other” background class	–	0.707
S9	3-class retrain (pre-water)	–	0.874
S10	Fine-tuning ablation (no-FT)	–	0.847
S11	Threshold grid, water-incl. final	–	0.704

D. Metrics

Two metrics on the **1,210**-frame own-test set (ride-disjoint from tuning; 983/40/307 veg/water/city positives). **Macro-F1**: per-class F1 (precision–recall harmonic mean) averaged equally over classes, so rare *water* counts fully. **Exact-match**: share of frames whose entire three-class label is correct at once. Together they expose each model’s strengths and decide the deployment.

E. Results and Decision

Results. SegFormer wins overall: macro-F1 0.704 vs. 0.649 and exact-match 0.755 vs. 0.691. The models tie on the common classes (vegetation 0.930/0.941, city 0.799/0.802), so the rare *water* class (3% of frames) decides it: SegFormer F1 0.382 vs. 0.206. Figure 4 localises the decision in the bottom-left – vegetation and city cluster on one high-F1 band, only the water pair is far apart – and shows that SegFormer wins water by *precision*, *not recall*: CLIP reaches higher recall (0.53 vs. 0.33) only by firing indiscriminately at precision 0.13, whereas SegFormer holds 0.46 for a higher F1.

Decision. We deploy the fine-tuned SegFormer: it wins the deciding class with trustworthy precision, ties CLIP elsewhere, is $40\times$ smaller (3.7M vs. 151M), keeping deployment feasible, and feeds the score area fractions directly – which CLIP cannot. *Water stays the known limitation*: recall ($\sim 67\%$ of true-water frames missed) is capped by data scarcity, fixable only with more in-domain water.




Figure 4. Precision–recall operating points with iso-F1 contours (own-test, 1,210 frames); arrows run CLIP $\rightarrow$ SegFormer. The isolated water pair (bottom-left) decides it: SegFormer trades recall for precision to reach a higher F1.

VII. COMPUTATION OF THE BIKEABILITY SCORE

The bikeability score fuses the three detector outputs into a single value in $[0, 100]$: the one-hot surface vector $\mathbf{g}$, environment area fractions $\mathbf{f}$, and object counts $\mathbf{n}$ are each mapped to a *sub-score* in $[0, 1]$ via a stored value vector (Table I), then combined by a weighted mean.

A. Overall Formula

For each evaluated frame, the score A results from the weighted mean of the three sub-scores:

$$A = 100 \cdot (w_g S_{\text{ground}} + w_e S_{\text{env}} + w_o S_{\text{object}}), \quad (1)$$

with the weights

$$w_g = 0.20, \quad w_e = 0.45, \quad w_o = 0.35. \quad (2)$$

The weights are heuristic: environment and object exposure dominate how a stretch feels and take the larger share ($w_e + w_o = 0.80$), while surface — which varies little and is the least reliable detector (macro-F1 0.58) — takes the smallest ($w_g = 0.20$). As all sub-scores lie in $[0, 1]$ and the weights sum to 1, $A \in [0, 100]$.

B. Sub-Scores

Each detector output maps to a sub-score in $[0, 1]$ via the value vectors of Table I, summarised in Table V.

Table V
PER-FRAME SUB-SCORE FORMULAS. THE DETECTOR OUTPUTS ARE $\mathbf{g}$ (ONE-HOT SURFACE), $\mathbf{f}$ (ENVIRONMENT AREA FRACTIONS) AND $\mathbf{n}$ (OBJECT COUNTS); $\mathbf{q}^g, \mathbf{q}^e, \boldsymbol{\lambda}$ ARE THE VALUE VECTORS OF TABLE I AND $\rho = \langle \mathbf{n}, \boldsymbol{\lambda} \rangle$.

Sub-score	Formula
Ground	$S_{\text{ground}} = \langle \mathbf{g}, \mathbf{q}^g \rangle$
Environment	$S_{\text{env}} = \frac{\langle \mathbf{f}, \mathbf{q}^e \rangle}{\sum_k f_k}$ (or 0.5 if $\sum_k f_k = 0$)
Object	$S_{\text{object}} = \exp(-\max(\rho, 0))$

The formulas tie to the detector outputs: **Ground** (one-hot) selects the surface's quality; **Environment** is a fraction-weighted mean (normalised by $\sum_k f_k$, 0.5 if empty) rewarding scenery coverage; **Object** clamps ρ so cars (+1.0) and traffic lights (+0.3) lower the score, bicycles (−0.2) only offset them, never lifting it above 1.0.




Figure 5. Bikeability score along the GPX route (red = low, green = high; mean $\bar{A} \approx 85.7/100$).

C. GPX Mapping and Visualization

Each frame is matched to its nearest GPX track point (nearest-neighbor join, 3 s tolerance, one manual sync point), giving every score a location. The route average $\bar{A} = \frac{1}{M} \sum_i A_i$ then colors short segments red-to-green, as mapped in Fig. 5: the example ride averages $\bar{A} \approx 85.7/100$, mostly green with a few yellow/orange sections.

VIII. CONCLUSION

We presented a pipeline that turns forward-facing cycling video into a bikeability score in $[0, 100]$ from three detectors: *surface* (training-free OpenCLIP ViT-B/32, macro-F1 0.58), *objects* (YOLO11s, beating the newer YOLO26 at equal scale), and *environment* (fine-tuned SegFormer over zero-shot CLIP, 0.704 vs. 0.649, winning the rare *water* class on precision and feeding the score’s area fractions). Water recall (scarce in-domain data) is the main limit; the next step is pixel-masked in-domain water frames and validation against human ratings.

REFERENCES

- [1] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2021, pp. 8748–8763.
- [2] G. Ilharco, M. Wortsman, R. Wightman, C. Gordon, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt, “OpenCLIP,” Zenodo, 2021, https://github.com/mlfoundations/open_clip.
- [3] S. Y. Gadre, G. Ilharco, A. Fang *et al.*, “DataComp: In search of the next generation of multimodal datasets,” in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [4] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, “Sigmoid loss for language image pre-training,” in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2023, pp. 11 975–11 986.
- [5] P. K. A. Vasu, H. Pouransari, F. Faghri, R. Vemulapalli, and O. Tuzel, “MobileCLIP: Fast image-text models through multi-modal reinforced training,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [6] Q. Sun, Y. Fang, L. Wu, X. Wang, and Y. Cao, “EVA-CLIP: Improved training techniques for CLIP at scale,” *arXiv preprint arXiv:2303.15389*, 2023.
- [7] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft COCO: Common objects in context,” in *Computer Vision – ECCV 2014*, 2014, pp. 740–755.
- [8] G. Jocher, J. Qiu, M. Liu, S. Lyu, F. C. Akyon, and M. E. Kalfaoglu, “Ultralytics YOLO26: Unified real-time end-to-end vision models,” *arXiv preprint arXiv:2606.03748*, 2026.
- [9] P. Martínez Otero, A. Tellaeché, and M. Hernández Melero, “Performance analysis of the YOLO object detection algorithm in embedded systems: Generated code vs. native implementation,” *Computation*, vol. 14, no. 3, p. 67, Mar. 2026.
- [10] E. Xie, W. Wang, Z. Yu, A. Anandkumar, J. M. Alvarez, and P. Luo, “SegFormer: Simple and efficient design for semantic segmentation with transformers,” in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2021, pp. 12 077–12 090.
- [11] G. Neuhold, T. Ollmann, S. Rota Bulò, and P. Kotschieder, “The Mapillary Vistas dataset for semantic understanding of street scenes,” in *Proc. IEEE Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 5000–5009.
- [12] B. Zhou, H. Zhao, X. Puig, T. Xiao, S. Fidler, A. Barriuso, and A. Torralba, “Semantic understanding of scenes through the ADE20K dataset,” *Int. J. Computer Vision (IJCV)*, vol. 127, no. 3, pp. 302–321, 2019.
- [13] T. Cornille and N. Rogge, “Fine-tune a semantic segmentation model with a custom dataset,” Hugging Face Blog, 2022, <https://huggingface.co/blog/fine-tune-segformer>.