

A Computer-Vision-Based Bikeability Score

Bünyamin Aydemir, Gonzalo Guerrero Salazar, and Hendrick Fischer

Abstract—

Index Terms—Bikeability, computer vision, zero-shot classification, vision-language models, CLIP, surface classification, object detection, GPX.

I. INTRODUCTION

Computer vision is one of the technologies that has firmly established itself in recent decades. The increase in computing capacity, the abundance of data, and the improvement of algorithms have propelled it—just as they have other areas of artificial intelligence—to the point where the technology is becoming a standard tool. We now view it as the de-facto solution to problems that previously required entirely different methods.

Take ADAS, for example, which are now expected in passenger vehicles rather than viewed as luxury equipment. However, these systems have largely not made their way to smaller vehicles like motorcycles or bicycles, the latter of which provide transportation and recreation to millions of people. Regular bicycle commuting inspired this project to implement vision systems not only to assess the safety, but also the aesthetic—and therefore mood-improving—aspects of bicycle routes. We call this metric the “bikeability score”, based on the premise that while all roads are technically drivable, they do not all provide the same riding experience.

The score singles out three distinct factors that define this experience:

- **The road itself (Surface):** Determining the type of surface the user is riding on to assess if the vehicle is adequate and the ride will be physically comfortable.
- **The obstacles along the route (Traffic Dynamics):** Evaluating whether traffic is dominated by other bicycles or cars, and identifying major crossings to determine if the ride can remain largely continuous.
- **The surroundings (Environment):** Recognizing whether the route is dominated by cluttered cityscapes or greenery, and detecting natural features like forests and rivers to assess if the route is enjoyable rather than just a functional commute.

These three factors are defined as our target sub-functions, as each contributes differently to the riding experience and operates under distinct dynamics. Consequently, each task required a dedicated approach, making it essential to select models specifically suited to each sub-function. Similarly, dataset annotation and labeling were tailored to the specific demands of each individual task. The following sections explain the methodology for selecting, analyzing, and adapting

these individual models. The final section details how the overall score is calculated from the combined outputs of each pipeline.

II. RESEARCH APPROACH AND METHODOLOGY

The system is implemented as a multi-stream pipeline. Video data, collected using a handlebar-mounted DJI camera, is sampled at a frequency of one frame every five seconds. This dataset is passed in batches of eight to the individual detector functions. The resulting feature arrays are then processed by the primary scoring function, with each frame’s evaluated score appended to an output file representing the main product of this analysis. Downstream from this output, several interpretation methods are possible, including min/max analysis and score averaging. For this project, a visualization was also generated by integrating GPS data.

The methodology initially leveraged pre-trained and zero-shot models where possible to minimize manual labeling effort and allow the taxonomy to evolve. However, as demonstrated in subsequent sections, this approach encountered specific limitations. As a result, each detector was independently benchmarked before integration into the final pipeline.

III. OBJECTIVES AND CONSTRAINTS

The primary objective is to demonstrate a functional pipeline capable of evaluating real-world data. Consequently, testing the framework on self-collected footage was a strict requirement. However, this imposed a key constraint: ground-truth labeling had to be generated internally, either manually or with semi-automated assistance. This limitation constrained both the volume of real-world ground truth data available and the scope of benchmarking that could be conducted on candidate models.

IV. ROAD-SURFACE CLASSIFICATION

The ridden surface strongly affects cycling comfort and safety: a paved cycleway is effortless, whereas loose gravel or bare soil is strenuous and, on a road bike, hazardous. Ground detection assigns each sampled frame a single surface class and supplies the surface term of the score. Frames come from a handlebar-mounted camera at one frame every five seconds; as each is dominated by one surface, the task is single-label classification, not segmentation.

A. Task Definition and Model Selection

Two model families can solve this task: a classifier *trained on our own labels*, or a *zero-shot* vision-language model (VLM). We avoid a self-trained classifier, which would need a large, balanced, manually labeled training set and re-training on every change of the class definitions. A zero-shot VLM of the CLIP family [1] instead classifies an image by its similarity to natural-language class descriptions, so the taxonomy is adapted by editing text alone.

B. Surface Classes and Prompts

We distinguish four cycling-relevant classes: **Cycleway** (a dedicated paved bike path), **Road** (a paved road for motor traffic), **Gravel** (loose crushed stone), and **Unpaved** (soil, mud, or field tracks). Each class is described by several English prompts whose normalized embeddings are averaged (*prompt ensembling*). As Cycleway and Road share the same smooth asphalt, the prompts encode the *discriminating context* (a narrow separated bicycle lane versus a wide multi-lane road) to separate the two embeddings. A frame is assigned to the class of highest cosine similarity.

C. Evaluation Protocol

We compare nine open CLIP-family models of light to medium size through the unified `open_clip` API [2]: OpenAI CLIP, OpenCLIP (LAION-2B, DataComp-XL) [3], SigLIP [4], MobileCLIP [5], and EVA-CLIP [6]; both *accuracy* and *latency* matter for on-bike use. For the quantitative evaluation we manually labeled more than 1,700 frames from self-recorded rides, excluding frames without a clear surface. As the classes are imbalanced, we report accuracy, balanced accuracy and macro-F1 with 95% bootstrap confidence intervals and McNemar tests; latency is measured at batch size 1 with warm-up runs (ms/frame and FPS).

D. Results

Table I lists the outcome, sorted by macro-F1. The best model, **OpenCLIP ViT-B/32 (LAION-2B)**, reaches macro-F1 0.58 at only 115 ms/frame, the best accuracy-latency trade-off. Two points stand out. First, the highest plain *accuracy* (0.78, CLIP ViT-B/16) does *not* give the highest macro-F1: accuracy rewards the majority classes, whereas macro-F1 weights all four surfaces equally and is the more faithful measure on imbalanced data. Second, the 95% bootstrap accuracy intervals of the leading models overlap, so their differences are *not*

Table I
ZERO-SHOT SURFACE-CLASSIFICATION BENCHMARK ON THE SELF-LABELED TEST SET (> 1,700 FRAMES), SORTED BY MACRO-F1.
BACC: BALANCED ACCURACY; LATENCY AT BATCH SIZE 1.

Model	P[M]	Acc	BAcc	mF1	ms	FPS
OpenCLIP B/32 (LAION-2B)	151	0.75	0.72	0.58	115	8.7
SigLIP B/16 (WebLI)	203	0.75	0.62	0.57	259	3.9
CLIP B/16 (OpenAI)	150	0.78	0.62	0.57	261	3.8
MobileCLIP-S2	99	0.73	0.70	0.56	304	3.3
CLIP B/32 (OpenAI)	151	0.70	0.60	0.54	93	10.7
EVA02-B/16	150	0.76	0.62	0.54	304	3.3
OpenCLIP B/32 (DataComp-XL)	151	0.75	0.65	0.52	125	8.0
CLIP L/14 (OpenAI)	428	0.64	0.50	0.42	929	1.1

statistically significant; only CLIP ViT-L/14 (428 M) falls clearly below them; the largest model is both slowest and weakest (0.42 macro-F1), confirming that capacity alone does not help.

Broken down per class, the best model scores F1 0.80 (Cycleway) and 0.73 (Road) but only 0.45 (Gravel) and 0.33 (Unpaved); as macro-F1 averages the four equally, the two strong classes are pulled down by the two weak ones (Fig. 1). Crucially, the weakness is one of *precision*, not recall: Unpaved is still recalled well (0.77), yet its precision collapses to 0.21 because, with only 13 true frames, a few false positives bleeding in from the majority classes suffice to ruin it; Gravel (51 frames) behaves the same way, while the residual Cycleway/Road confusion (both smooth asphalt) is comparatively minor. A macro-F1 of 0.58 is thus by design strict on such tiny classes; the accuracy (0.75) and balanced accuracy (0.72, the highest of all models) are more flattering yet still fair, and for a zero-shot classifier on strongly imbalanced data the result is solid. We therefore adopt OpenCLIP ViT-B/32 (LAION-2B) as the ground detector; as all models share the same pipeline, the differences reflect architecture and training data (latencies are hardware-dependent).




Figure 1. Row-normalized confusion matrix of the best model (OpenCLIP ViT-B/32, LAION-2B) on the labeled test set. Each cell shows the count and the row share; the diagonal is correct. Errors concentrate on the Cycleway/Road pair and the rare Gravel/Unpaved classes.

V. OBJECT DETECTION

A. Technical Scope

For the object recognition module, the methodology focuses specifically on counting objects. This design choice was made because objects move relative to the bike rider regardless of their absolute speed, as the rider is in continuous motion. Consequently, an object's size within the field of view changes frame by frame and cannot serve as a stable factor for calculations.

This dynamic property led to another design consideration regarding the scope of the “bikeability score” evaluator: evaluating the exact same route will yield slightly different results depending on the specific time the footage is captured. For this algorithm to be viable in a mass-market product, it must be capable of running in real time on an edge device mounted on the bicycle. This strict real-time requirement applies to the dynamic object detection, but not to the other two static sub-scores.

B. Performance Trade-offs

Therefore, the primary selection parameters for this specific task are the model's inference time and its overall size, which are inherently related. Because larger models can extract more complex image features, there is a direct trade-off between inference time and accuracy, requiring us to find an optimal balance for our model selection.

These runtime criteria, combined with the fact that this module only requires object recognition, led directly to targeting the YOLO family of models as the most suitable candidates for the task. However, it remains necessary to benchmark different generations and sizes of YOLO models to determine the most effective configuration.




Figure 2. YOLO model comparison for the object-detection module. (a) Inference speed vs. macro MAE across YOLO variants. (b) Macro MAE vs. confidence threshold per model.

C. Object Classes

For the calculation of the score, the methodology restricts the categories of relevant objects. Rather than defining new categories from scratch, the project utilizes widely accepted object set categorizations. Because the system employs a pre-trained YOLO model, it is more efficient to leverage its baseline capabilities before resorting to custom retraining. Consequently, the COCO [7] dataset convention was selected as the primary starting point, though Mapillary Vistas [8] remains a viable alternative dataset for future iterations.

The COCO dataset classifies detected objects into 80 distinct categories, not all of which are relevant to traffic environments. Therefore, the relevant classes were grouped into three primary subsets (Table II).

Table II
GROUPING OF RELEVANT COCO OBJECT CLASSES INTO TRAFFIC-RELATED CATEGORIES.

Category	COCO Object Classes
Motor vehicles	car, truck, bus, motorcycle
Other road users	person, bicycle
Traffic infrastructure	traffic light, stop sign, parking meter, fire hydrant, bench

From these subsets, three representative classes were selected: **car**, **bicycle**, and **traffic light**.

To ensure accurate labeling within the reference dataset, it is necessary to strictly define the boundaries and potential edge cases for the target classes (Table III).

Table III
CLASS BOUNDARY DEFINITIONS FOR THE OBJECT DETECTOR.

Class	IS	IS NOT
Bicycle	Road bike, mountain bike, city bike, BMX, w/rider	Motorcycle, motor scooter
Car	Sedan, hatchback, wagon, coupe, SUV, crossover, minivan	Bus, truck, train

VI. ENVIRONMENT DETECTION

VII. COMPUTATION OF THE BIKEABILITY SCORE

The bikeability score combines the three components (surface from ground-detection, environment from environment-detection, and objects from object-detection) into a single value in the interval $[0, 100]$. Every detector emits a fixed-length *vector*; each vector is first mapped to a *sub-score* in $[0, 1]$, and the three sub-scores are then combined by a weighted mean.

A. Detector Outputs as Vectors

The three detectors deliberately use different but fixed output representations, each aligned element by element with a stored value vector:

- **Ground:** a one-hot vector $\mathbf{g} \in \{0, 1\}^4$ over the classes (Cycleway, Road, Gravel, Unpaved); exactly one entry is 1.
- **Environment:** a vector of area fractions $\mathbf{f} \in [0, 1]^3$ over (vegetation, water, city).
- **Object:** a vector of counts $\mathbf{n} \in \mathbb{Z}_{\geq 0}^3$ over (bicycle, car, traffic light).

B. Overall Formula

For each evaluated frame, the score A results from the weighted mean of the three sub-scores:

$$A = 100 \cdot (w_g S_{\text{ground}} + w_e S_{\text{env}} + w_o S_{\text{object}}), \quad (1)$$

with the weights

$$w_g = 0.20, \quad w_e = 0.45, \quad w_o = 0.35. \quad (2)$$

The weights are heuristic and follow two considerations. Environment and object exposure jointly capture how pleasant and how safe a stretch feels, the factors a cyclist perceives most directly, and therefore receive the larger share ($w_e + w_o = 0.80$). The surface receives the smallest weight ($w_g = 0.20$): on typical routes it varies little, and ground detection is also the least reliable component (macro-F1 0.58), so a smaller weight limits the influence of its errors. Since all sub-scores lie in $[0, 1]$ and the weights sum to 1, the score is bounded to $[0, 100]$. In practice the lower extreme 0 is never reached, as $S_{\text{ground}} \geq 0.1$ and $S_{\text{object}} = \exp(-\max(\rho, 0)) > 0$; the upper value 100, however, does occur whenever all three sub-scores equal 1 simultaneously, e.g. on a paved cycleway ($S_{\text{ground}} = 1$) through vegetation or along water ($S_{\text{env}} = 1$) with no motorised traffic ($S_{\text{object}} = 1$).

C. Ground Sub-Score

The ground sub-score is the inner product of the one-hot vector $\mathbf{g}$ with the surface quality vector $\mathbf{q}^g = (1.0, 0.7, 0.3, 0.1)$ (Table IV):

$$S_{\text{ground}} = \langle \mathbf{g}, \mathbf{q}^g \rangle = \sum_k g_k q_k^g. \quad (3)$$

Because $\mathbf{g}$ is one-hot, this simply selects the quality of the detected surface class.

D. Environment Sub-Score

The environment sub-score is the fraction-weighted mean of the environment quality vector $\mathbf{q}^e = (1.0, 1.0, 0.4)$:

$$S_{\text{env}} = \begin{cases} \frac{\langle \mathbf{f}, \mathbf{q}^e \rangle}{\sum_k f_k}, & \text{if } \sum_k f_k > 0, \\ 0.5, & \text{otherwise.} \end{cases} \quad (4)$$

The normalization by $\sum_k f_k$ keeps the sub-score in $[0, 1]$ even when the detected fractions do not cover the whole frame; a neutral value of 0.5 is used if no environment class is detected.

E. Object Sub-Score

The object sub-score maps the counts $\mathbf{n}$ through a *saturation function* using the penalty vector $\boldsymbol{\lambda} = (-0.2, 1.0, 0.3)$:

$$S_{\text{object}} = \exp\left(-\max(\rho, 0)\right), \quad \rho = \langle \mathbf{n}, \boldsymbol{\lambda} \rangle, \quad (5)$$

where ρ is evaluated independently per frame. A positive factor for cars (+1.0) and traffic lights (+0.3) makes ρ positive and thus lowers the sub-score, whereas bicycles carry a *negative* factor (−0.2). The clamp $\max(\rho, 0)$ makes the score asymmetric: a frame is never rewarded *above* the neutral value 1.0, so bicycles do not inflate the score on their own; they only *offset* the penalty of nearby cars or traffic lights. This encodes the intuition that good cycling conditions are the default and only motorised traffic degrades them. The exponential form additionally caps the influence of outliers, so a single crowded frame cannot dominate the overall score.

Table IV
CLASS VALUE VECTORS FOR THE THREE SUB-SCORES: SURFACE QUALITY $\mathbf{q}^g$, ENVIRONMENT QUALITY $\mathbf{q}^e$, AND OBJECT WEIGHT $\boldsymbol{\lambda}$ (HEURISTIC; A NEGATIVE VALUE ACTS AS A BONUS).

Component	Class	Value
Ground $\mathbf{q}^g$	Cycleway	1.0
	Road	0.7
	Gravel	0.3
	Unpaved	0.1
Environment $\mathbf{q}^e$	Vegetation	1.0
	Water	1.0
	City	0.4
Object weight $\boldsymbol{\lambda}$	Bicycle	−0.2
	Car	+1.0
	Traffic light	+0.3

F. GPX Mapping and Visualization

Because the camera clock is unreliable, the frames are aligned to the GPS track by a single *manual sync point* (by default: recording start = GPX track start). Each frame is then matched to its nearest GPX track point in time (a nearest-neighbor join within a 3 s tolerance: large enough to bridge gaps in the roughly 1 Hz GPS log, yet below the 5 s frame spacing so that each frame maps to a distinct track point), which assigns a latitude/longitude to every score. The *average* bikeability score of the route is the mean $\bar{A} = \frac{1}{M} \sum_{i=1}^M A_i$ over all M evaluated frames. Finally, the route is drawn as a sequence of short segments, each colored by the local score

on a red–orange–yellow–green scale (Fig. 3). For the example ride shown, the mean score is $\bar{A} \approx 85.7/100$: the track is predominantly green (good cycling conditions) with a few yellow/orange sections.




Figure 3. Bikeability score colored along the GPX route (red = low, green = high). Per-frame scores are matched to the nearest GPS track point via a manual sync point; the example ride has a mean score of $\bar{A} \approx 85.7/100$.

VIII. DISCUSSION

IX. CONCLUSION

REFERENCES

- [1] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *Proc. Int. Conf. Machine Learning (ICML)*, 2021, pp. 8748–8763.
- [2] G. Ilharco, M. Wortsman, R. Wightman, C. Gordon, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt, "OpenCLIP," Zenodo, 2021, https://github.com/mlfoundations/open_clip.
- [3] S. Y. Gadre, G. Ilharco, A. Fang *et al.*, "DataComp: In search of the next generation of multimodal datasets," in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [4] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, "Sigmoid loss for language image pre-training," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2023, pp. 11 975–11 986.
- [5] P. K. A. Vasu, H. Pouransari, F. Faghri, R. Vemulapalli, and O. Tuzel, "MobileCLIP: Fast image-text models through multi-modal reinforced training," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [6] Q. Sun, Y. Fang, L. Wu, X. Wang, and Y. Cao, "EVA-CLIP: Improved training techniques for CLIP at scale," *arXiv preprint arXiv:2303.15389*, 2023.
- [7] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, "Microsoft COCO: Common objects in context," in *Computer Vision – ECCV 2014*, 2014, pp. 740–755.
- [8] G. Neuhold, T. Ollmann, S. Rota Bulò, and P. Kotschieder, "The Mapillary Vistas dataset for semantic understanding of street scenes," in *Proc. IEEE Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 5000–5009.